

Using CNNs to classify people in Famous48 dataset - Subject 12c

Jakub Czernański
193105

Paweł Blicharz
193193

02.05.2024

Contents

Contents

1	Introduction	1
2	Libraries used	1
3	File loading	2
4	Testing	3
4.1	Architectures	3
4.1.1	AlexNet	3
4.1.2	LeNet-5	5
5	Final model	6

1 Introduction

Lorem ipsum

Data description

famous48 is a set of example images contained faces of 48 famous persons like sportsmen, politicians, actors or television stars. It was divided into 3 files: x24x24.txt, y24x24.txt, z24x24.txt, each containing 16 personal classes.

Attributes description:

- a1 - face containing flag: (1-with face, 0-without face)
- a2 - image number in current class (person) beginning from 0
- a3 - class (person) number beginning from 0
- a4 - sex (0 - woman, 1 - man)
- a5 - race (0- white, 1 - negro, 2 - indian, ...)
- a6 - age (0 - baby, 1 - young, 2 - middle-age, 3 - old)
- a7 - binokulars (0 - without, 1 - transparent, 2 - dark)
- a8 - emotional expression (not state!) (0 - sad, 1 - neutral, 2 - happy)

Note:

Full code can be found in *notebook_keras.ipynb*. Due to limited space, we provide here only the most important code.

2 Libraries used

Firstly, we import all our libraries:

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow import keras
4
5 from keras import layers
6 from keras import regularizers
7 from keras import backend as K
8 from keras import Sequential, Input
9 from keras.optimizers import SGD, Adam
10 from keras.losses import categorical_crossentropy
11 from keras.callbacks import LearningRateScheduler
12
13 from sklearn.model_selection import train_test_split
14
15 import matplotlib.pyplot as plt
16
17 import math
18 from functools import partial
```

3 File loading

This is our code for loading files and we will not change it throughout our journey:

```
1 CLASSES = 48
2 IMAGE_SIZE = 24
3
4 def read_file(filename):
5
6     with open(filename, 'r') as file:
7         file.readline() # we skip the first line as it is not needed
8         number_of_pixels = int(file.readline())
9
10        features = []
11        labels = []
12
13        for line in file.readlines():
14            elements = line.split()
15
16            # add features
17            pixels = np.array(elements[:number_of_pixels], dtype=float)
18            pixels = np.reshape(pixels, [IMAGE_SIZE, IMAGE_SIZE])
19            features.append(pixels)
20
21            # add labels
22            labels.append(elements[number_of_pixels+2])
23
24        features = np.array(features)
25        labels = np.array(labels, dtype=int)
26
27        return features, labels
28
29
30 X_0, y_0 = read_file('./data/x24x24.txt')
31 X_1, y_1 = read_file('./data/y24x24.txt')
32 X_2, y_2 = read_file('./data/z24x24.txt')
33
34 X = np.concatenate((X_0, X_1, X_2))
35 y = np.concatenate((y_0, y_1, y_2))
36
37 X_train, X_test, y_train_raw, y_test_raw = train_test_split(X, y,
38                                                                train_size=N_TRAIN_EXAMPLES,
39                                                                test_size=N_TEST_EXAMPLES,
40                                                                random_state=42)
41
42 y_train = np.zeros((y_train_raw.shape[0], CLASSES))
43 y_test = np.zeros((y_test_raw.shape[0], CLASSES))
```

```

44
45 for i, value in enumerate(y_train_raw):
46     y_train[i][value] = 1
47
48 for i, value in enumerate(y_test_raw):
49     y_test[i][value] = 1

```

4 Testing

Here are our all attempts at finding the best model. We tried out different architectures and hyperparameters to achieve that goal.

4.1 Architectures

4.1.1 AlexNet

Firstly, we implemented AlexNet model (https://papers.nips.cc/paper_files/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html) which won the 2012 ILFRC challenge. Here is our model in python code, after downscaling it and changing some parameters:

```

1 conv_regularizer = regularizers.l2(0.0006)
2 dense_regularizer = regularizers.l2(0.01)
3
4 DefaultConv2D = partial(tf.keras.layers.Conv2D, kernel_size=3, padding="same",
5                         activation="relu", kernel_regularizer=conv_regularizer)
6
7 model = keras.Sequential([
8     Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1)),
9     DefaultConv2D(96),
10    layers.MaxPooling2D(pool_size=3, strides=2),
11
12    tf.keras.layers.Dropout(0.3),
13    DefaultConv2D(256, kernel_size=5),
14    tf.keras.layers.MaxPooling2D(pool_size=3, strides=2),
15
16    tf.keras.layers.Dropout(0.4),
17    DefaultConv2D(384),
18    tf.keras.layers.Dropout(0.5),
19    DefaultConv2D(384),
20    tf.keras.layers.MaxPooling2D(pool_size=3, strides=2),
21
22    tf.keras.layers.Flatten(),
23    tf.keras.layers.Dropout(0.6),
24    tf.keras.layers.Dense(384, activation='relu',
25                          kernel_regularizer=dense_regularizer),

```

```

26     tf.keras.layers.Dense(CLASSES, activation='softmax')
27 ])

```

Hyperparameters & methods used:

- *optimizer*: Adam, with *learning rate* = 0.001
- *epochs* = 250
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers: $l = 0.0006$
- L2 regularization for the dense layer: $l = 0.01$

Results: (rounded to 3 decimal places)

- train set accuracy: 0.896
- train loss: 0.775
- validation set accuracy: 0.81
- validation set loss: 1.115
- test set accuracy: **0.8**
- test set loss: 1.137

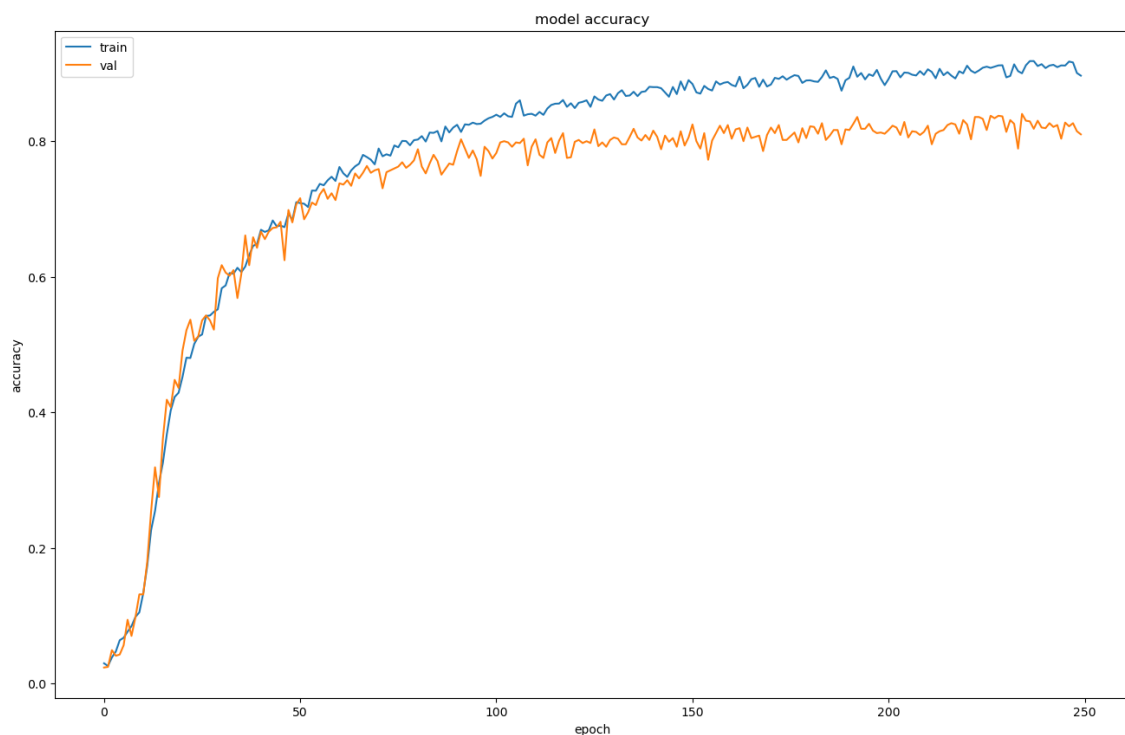


Figure 1: Model accuracy for AlexNet architecture - training & validation sets

Conclusions:

We decided that 80% is not enough so we did not tune hyperparameters.

4.1.2 LeNet-5

Next, we decided to use LeNet-5 architecture, presented below:

```
1 conv_regularizer = regularizers.l2(0.0009096443481619992)
2 dense_regularizer = regularizers.l2(0.011905583599301073)
3
4 dropout_base = 0.09439855997376015
5 dropout_inc = 0.14131761625994724
6 dropout_1 = dropout_base
7 dropout_2 = dropout_base + dropout_inc
8 dropout_3 = dropout_base + 2*dropout_inc
9
10 DefaultConv2D = partial(tf.keras.layers.Conv2D, kernel_size=5,
11                        padding="same", activation="tanh",
12                        kernel_regularizer=conv_regularizer)
13
14 model = Sequential([
15     Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1)),
16     DefaultConv2D(6),
17     layers.MaxPooling2D(pool_size=2, strides=2),
18
19     layers.Dropout(dropout_1),
20     DefaultConv2D(16),
21     layers.MaxPooling2D(pool_size=2, strides=2),
22
23     layers.Dropout(dropout_2),
24     DefaultConv2D(120),
25
26     layers.Flatten(),
27     layers.Dropout(dropout_3),
28     DefaultConv2D(84),
29     layers.Dense(CLASSES, activation='softmax')
30 ])
```

Hyperparameters & methods used:

- *optimizer*: Adam, with *learning rate* = 0.001
- *epochs* = 100
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers: $l \approx 0.00091$
- L2 regularization for the dense layer: $l \approx ?(to - do)$

Results: (rounded to 2(to-do: change to 3) decimal places)

- train set accuracy: 0.92
- train loss: ?
- validation set accuracy: 0.82
- validation set loss: ?
- test set accuracy: **0.82**
- test set loss: ?




Figure 2: Model accuracy for LeNet-5 architecture - training & validation sets

Conclusions:

The model after 20 epochs starts overfitting and validation set's accuracy improves very slowly.

5 Final model