

# Using CNNs to classify people in Famous48 dataset - Subject 12c

Jakub Czernański  
193105

Paweł Blicharz  
193193

02.05.2024

# Contents

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Libraries used</b>	<b>1</b>
<b>3</b>	<b>File loading</b>	<b>2</b>
<b>4</b>	<b>Testing</b>	<b>3</b>
4.1	AlexNet . . . . .	3
4.2	LeNet-5 . . . . .	5
4.2.1	Base model . . . . .	5
4.2.2	Testing hyperparameters . . . . .	5
<b>5</b>	<b>Summary</b>	<b>8</b>

# 1 Introduction

In this project we have 24x24 images of famous people and our goal is to create a model that can predict who the person is given the image. In order to do it, we divided our data into training, validation and test sets (validation set is created during learning - in `fit()` function). We use Convolutional Neural Networks and tested 2 popular CNN architectures: AlexNet and LeNet-5.

## Data description

famous48 is a set of example images contained faces of 48 famous persons like sportsmens, politicians, actors or television stars. It was divided into 3 files: `x24x24.txt`, `y24x24.txt`, `z24x24.txt`, each containing 16 personal classes.

Attributes description:

- a1 - face containing flag: (1-with face, 0-without face)
- a2 - image number in current class (person) beginning from 0
- a3 - class (person) number beginning from 0
- a4 - sex (0 - woman, 1 - man)
- a5 - race (0- white, 1 - negro, 2 - indian, ...)
- a6 - age (0 - baby, 1 - young, 2 - middle-age, 3 - old)
- a7 - binokulars (0 - without, 1 - transparent, 2 - dark)
- a8 - emotional expression (not state!) (0 - sad, 1 - neutral, 2 - happy)

**Note:**

Full code can be found in *notebook\_keras.ipynb*, *lenet.ipynb*, *alexnet.ipynb* and *after\_optuna\_training.ipynb*. Due to limited space, we provide here only the most important code.

## 2 Libraries used

Firstly, we import all our libraries:

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow import keras
4
5 from keras import layers
6 from keras import regularizers
7 from keras import backend as K
8 from keras import Sequential, Input
9 from keras.optimizers import SGD, Adam
10 from keras.losses import categorical_crossentropy
11 from keras.callbacks import LearningRateScheduler
12
13 from sklearn.model_selection import train_test_split
14
15 import matplotlib.pyplot as plt
16
```

```

17 import math
18 from functools import partial
19
20 from utils.constants import CLASSES, IMAGE_SIZE
21 from utils.load_dataset import load_dataset

```

### 3 File loading

This is our code for loading files and we will not change it throughout our journey:

```

1 CLASSES = 48
2 IMAGE_SIZE = 24
3
4 def read_file(filename):
5     with open(filename, 'r') as file:
6         file.readline() # we skip the first line as it is not needed
7         number_of_pixels = int(file.readline())
8
9         features = []
10        labels = []
11
12        for line in file.readlines():
13            elements = line.split()
14
15            # add features
16            pixels = np.array(elements[:number_of_pixels], dtype=float)
17            pixels = np.reshape(pixels, [IMAGE_SIZE, IMAGE_SIZE])
18            features.append(pixels)
19
20            # add labels
21            labels.append(elements[number_of_pixels+2])
22
23        features = np.array(features)
24        labels = np.array(labels, dtype=int)
25
26    return features, labels
27
28
29 def load_dataset(directory_path):
30     X_0, y_0 = read_file(f'{directory_path}/x24x24.txt')
31     X_1, y_1 = read_file(f'{directory_path}/y24x24.txt')
32     X_2, y_2 = read_file(f'{directory_path}/z24x24.txt')
33
34     X = np.concatenate((X_0, X_1, X_2))
35     y = np.concatenate((y_0, y_1, y_2))
36
37     N_TRAIN_EXAMPLES=int(len(X) * 0.8)
38     N_TEST_EXAMPLES=len(X) - N_TRAIN_EXAMPLES
39
40     X_train, X_test, y_train_raw, y_test_raw = train_test_split(X, y,

```

```

41                                     train_size=N_TRAIN_EXAMPLES,
42                                     test_size=N_TEST_EXAMPLES,
43                                     random_state=42)
44
45     y_train = np.zeros((y_train_raw.shape[0], CLASSES))
46     y_test = np.zeros((y_test_raw.shape[0], CLASSES))
47
48     for i, value in enumerate(y_train_raw):
49         y_train[i][value] = 1
50
51     for i, value in enumerate(y_test_raw):
52         y_test[i][value] = 1
53
54     return X_train, X_test, y_train, y_test
55
56
57 DIRPATH = './data'
58 X_train, X_test, y_train, y_test = load_dataset(DIRPATH)

```

## 4 Testing

Here are our all attempts at finding the best model. We tried out different architectures and hyperparameters to achieve that goal.

### 4.1 AlexNet

Firstly, we implemented AlexNet model ([https://papers.nips.cc/paper\\_files/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html](https://papers.nips.cc/paper_files/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html)) which won the 2012 ILFRC challenge. Here is our model in python code, after downscaling it and changing some parameters:

```

1  conv_regularizer = regularizers.l2(0.0006)
2  dense_regularizer = regularizers.l2(0.01)
3
4  DefaultConv2D = partial(tf.keras.layers.Conv2D, kernel_size=3, padding="same",
5                          activation="relu", kernel_regularizer=conv_regularizer)
6
7  model = keras.Sequential([
8      Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1)),
9      DefaultConv2D(96),
10     layers.MaxPooling2D(pool_size=3, strides=2),
11
12     tf.keras.layers.Dropout(0.3),
13     DefaultConv2D(256, kernel_size=5),
14     tf.keras.layers.MaxPooling2D(pool_size=3, strides=2),
15
16     tf.keras.layers.Dropout(0.4),
17     DefaultConv2D(384),
18     tf.keras.layers.Dropout(0.5),
19     DefaultConv2D(384),
20     tf.keras.layers.MaxPooling2D(pool_size=3, strides=2),

```

```

21
22     tf.keras.layers.Flatten(),
23     tf.keras.layers.Dropout(0.6),
24     tf.keras.layers.Dense(384, activation='relu',
25                             kernel_regularizer=dense_regularizer),
26     tf.keras.layers.Dense(CLASSES, activation='softmax')
27 ])
```

### Hyperparameters & methods used:

- *optimizer*: Adam, with *learning rate* = 0.001
- *epochs* = 250
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers:  $l = 0.0006$
- L2 regularization for the dense layer:  $l = 0.01$

### Results: (rounded to 3 decimal places)

- train set accuracy: 0.896
- train loss: 0.775
- validation set accuracy: 0.81
- validation set loss: 1.115
- test set accuracy: **0.8**
- test set loss: 1.137




Figure 1: Model accuracy for AlexNet architecture - training & validation sets

## Conclusions

We decided that 80% is not enough so we did not tune hyperparameters.

## 4.2 LeNet-5

Next, we decided to use LeNet-5 architecture. a

### 4.2.1 Base model

```
1 conv_regularizer = regularizers.l2(0.0006)
2 dense_regularizer = regularizers.l2(0.01)
3
4 DefaultConv2D = partial(tf.keras.layers.Conv2D, kernel_size=5, padding="same",
5                         activation="tanh", kernel_regularizer=conv_regularizer)
6
7 model = Sequential(
8     [
9         Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1)),
10        DefaultConv2D(6),
11        layers.MaxPooling2D(pool_size=2, strides=2),
12
13        layers.Dropout(0.45),
14        DefaultConv2D(16),
15        layers.MaxPooling2D(pool_size=2, strides=2),
16
17        layers.Dropout(0.3),
18        DefaultConv2D(120),
19
20        layers.Flatten(),
21        layers.Dropout(0.15),
22        layers.Dense(84, activation=activation_def,
23                    kernel_regularizer=dense_regularizer),
24        layers.Dense(CLASSES, activation='softmax'),
25    ]
26 )
```

### 4.2.2 Testing hyperparameters

#### Attempt 1

We decided to stay with Adam optimizer and *learning rate* = 0.001. We select regularization hyperparameters by hand (0.0006 and 0.01).

#### Hyperparameters & methods used:

- *optimizer*: Adam, with *learning rate* = 0.001
- *epochs* = 100
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers:  $l = 0.0006$
- L2 regularization for the dense layer:  $l = 0.01$

## Results:

- train set accuracy: 0.805
- train loss: 1.07
- validation set accuracy: 0.81
- validation set loss: 1.08
- test set accuracy: **0.794**
- test set loss: 1.095




Figure 2: Model accuracy for LeNet-5 architecture - training & validation sets, 100 epochs

## Conclusions

Results looked promising so we decided to test this model for 1000 epochs but unfortunately it started overfitting pretty soon and we achieved only **84.6%** test set accuracy. Here is a graph representing results for the same model but for 1000 epochs:






Figure 3: Model accuracy for LeNet-5 architecture - training & validation sets, 1000 epochs

## Attempt 2

We decided to find better hyperparameters ( $l$ ) for regularization and we used optuna to do this (the values are presented below).

We also introduced decaying learning rate and trained the model for 1000 epochs. This time we got much better results.

### Hyperparameters & methods used:

- *optimizer*: Adam, with starting *learning rate* = 0.001, and decreasing by 50% each 150 epochs
- *epochs* = 1000
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers:  $l \approx 0.00061$
- L2 regularization for the dense layer:  $l \approx 0.0301$

### Results:

- train set accuracy: 0.981
- train loss: 0.337
- validation set accuracy: 0.878
- validation set loss: 0.651
- test set accuracy: **0.892**

- test set loss: 0.629




Figure 4: Model accuracy for LeNet-5 with optimized hyperparameters

## Conclusions

Hyperparameter tuning along with decaying learning rate helped us achieve higher prediction accuracy - from **84.6%** in our first attempt to **89.2%** in the second attempt. We are satisfied with the resulting model and its results.

## 5 Summary

To sum up, we used 2 different CNN architectures to train our models and while AlexNet performed pretty poorly, the LeNet, after some hyperparameter tuning, achieved very good results: **89.2%** test set accuracy.