

Using CNNs to classify people in Famous48 dataset - Subject 12c

Jakub Czernański
193105

Paweł Blicharz
193193

02.05.2024

Contents

Contents

1	Introduction	1
2	Libraries used	1
3	File loading	2
4	Testing	3
4.1	AlexNet	3
4.2	LeNet-5	5
4.2.1	Base model	5
4.2.2	Testing hyperparameters	6
5	Final model	8

1 Introduction

Lorem ipsum

Data description

famous48 is a set of example images contained faces of 48 famous persons like sportsmen, politicians, actors or television stars. It was divided into 3 files: x24x24.txt, y24x24.txt, z24x24.txt, each containing 16 personal classes.

Attributes description:

- a1 - face containing flag: (1-with face, 0-without face)
- a2 - image number in current class (person) beginning from 0
- a3 - class (person) number beginning from 0
- a4 - sex (0 - woman, 1 - man)
- a5 - race (0- white, 1 - negro, 2 - indian, ...)
- a6 - age (0 - baby, 1 - young, 2 - middle-age, 3 - old)
- a7 - binokulars (0 - without, 1 - transparent, 2 - dark)
- a8 - emotional expression (not state!) (0 - sad, 1 - neutral, 2 - happy)

Note:

Full code can be found in *notebook_keras.ipynb*. Due to limited space, we provide here only the most important code.

2 Libraries used

Firstly, we import all our libraries:

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow import keras
4
5 from keras import layers
6 from keras import regularizers
7 from keras import backend as K
8 from keras import Sequential, Input
9 from keras.optimizers import SGD, Adam
10 from keras.losses import categorical_crossentropy
11 from keras.callbacks import LearningRateScheduler
12
13 from sklearn.model_selection import train_test_split
14
15 import matplotlib.pyplot as plt
16
17 import math
```

```

18 from functools import partial
19
20 from utils.constants import CLASSES, IMAGE_SIZE
21 from utils.load_dataset import load_dataset

```

3 File loading

This is our code for loading files and we will not change it throughout our journey:

```

1 CLASSES = 48
2 IMAGE_SIZE = 24
3
4 def read_file(filename):
5     with open(filename, 'r') as file:
6         file.readline() # we skip the first line as it is not needed
7         number_of_pixels = int(file.readline())
8
9         features = []
10        labels = []
11
12        for line in file.readlines():
13            elements = line.split()
14
15            # add features
16            pixels = np.array(elements[:number_of_pixels], dtype=float)
17            pixels = np.reshape(pixels, [IMAGE_SIZE, IMAGE_SIZE])
18            features.append(pixels)
19
20            # add labels
21            labels.append(elements[number_of_pixels+2])
22
23        features = np.array(features)
24        labels = np.array(labels, dtype=int)
25
26        return features, labels
27
28
29 def load_dataset(directory_path):
30     X_0, y_0 = read_file(f'{directory_path}/x24x24.txt')
31     X_1, y_1 = read_file(f'{directory_path}/y24x24.txt')
32     X_2, y_2 = read_file(f'{directory_path}/z24x24.txt')
33
34     X = np.concatenate((X_0, X_1, X_2))
35     y = np.concatenate((y_0, y_1, y_2))
36
37     N_TRAIN_EXAMPLES=int(len(X) * 0.8)

```

```

38     N_TEST_EXAMPLES=len(X) - N_TRAIN_EXAMPLES
39
40     X_train, X_test, y_train_raw, y_test_raw = train_test_split(X, y,
41                                                                 train_size=N_TRAIN_EXAMPLES,
42                                                                 test_size=N_TEST_EXAMPLES,
43                                                                 random_state=42)
44
45     y_train = np.zeros((y_train_raw.shape[0], CLASSES))
46     y_test = np.zeros((y_test_raw.shape[0], CLASSES))
47
48     for i, value in enumerate(y_train_raw):
49         y_train[i][value] = 1
50
51     for i, value in enumerate(y_test_raw):
52         y_test[i][value] = 1
53
54     return X_train, X_test, y_train, y_test
55
56
57 DIRPATH = './data'
58 X_train, X_test, y_train, y_test = load_dataset(DIRPATH)

```

4 Testing

Here are our all attempts at finding the best model. We tried out different architectures and hyperparameters to achieve that goal.

4.1 AlexNet

Firstly, we implemented AlexNet model (https://papers.nips.cc/paper_files/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html) which won the 2012 ILFRC challenge. Here is our model in python code, after downscaling it and changing some parameters:

```

1 conv_regularizer = regularizers.l2(0.0006)
2 dense_regularizer = regularizers.l2(0.01)
3
4 DefaultConv2D = partial(tf.keras.layers.Conv2D, kernel_size=3, padding="same",
5                        activation="relu", kernel_regularizer=conv_regularizer)
6
7 model = keras.Sequential([
8     Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1)),
9     DefaultConv2D(96),
10    layers.MaxPooling2D(pool_size=3, strides=2),
11
12    tf.keras.layers.Dropout(0.3),

```

```

13     DefaultConv2D(256, kernel_size=5),
14     tf.keras.layers.MaxPooling2D(pool_size=3, strides=2),
15
16     tf.keras.layers.Dropout(0.4),
17     DefaultConv2D(384),
18     tf.keras.layers.Dropout(0.5),
19     DefaultConv2D(384),
20     tf.keras.layers.MaxPooling2D(pool_size=3, strides=2),
21
22     tf.keras.layers.Flatten(),
23     tf.keras.layers.Dropout(0.6),
24     tf.keras.layers.Dense(384, activation='relu',
25                             kernel_regularizer=dense_regularizer),
26     tf.keras.layers.Dense(CLASSES, activation='softmax')
27 ])
```

Hyperparameters & methods used:

- *optimizer*: Adam, with *learning rate* = 0.001
- *epochs* = 250
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers: $l = 0.0006$
- L2 regularization for the dense layer: $l = 0.01$

Results: (rounded to 3 decimal places)

- train set accuracy: 0.896
- train loss: 0.775
- validation set accuracy: 0.81
- validation set loss: 1.115
- test set accuracy: **0.8**
- test set loss: 1.137




Figure 1: Model accuracy for AlexNet architecture - training & validation sets

Conclusions:

We decided that 80% is not enough so we did not tune hyperparameters.

4.2 LeNet-5

Next, we decided to used LeNet-5 architecture. a

4.2.1 Base model

```

1 conv_regularizer = regularizers.l2(0.0006)
2 dense_regularizer = regularizers.l2(0.01)
3
4 DefaultConv2D = partial(tf.keras.layers.Conv2D, kernel_size=5, padding="same",
5                           activation="tanh", kernel_regularizer=conv_regularizer)
6
7 model = Sequential(
8     [
9         Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1)),
10        DefaultConv2D(6),
11        layers.MaxPooling2D(pool_size=2, strides=2),
12
13        layers.Dropout(0.45),

```

```

14     DefaultConv2D(16),
15     layers.MaxPooling2D(pool_size=2, strides=2),
16
17     layers.Dropout(0.3),
18     DefaultConv2D(120),
19
20     layers.Flatten(),
21     layers.Dropout(0.15),
22     layers.Dense(84, activation=activation_def,
23                 kernel_regularizer=dense_regularizer),
24     layers.Dense(CLASSES, activation='softmax'),
25 ]
26 )

```

4.2.2 Testing hyperparameters

Attempt 1

We select all hyperparameters by hand

Hyperparameters & methods used:

- *optimizer*: Adam, with *learning rate* = 0.001
- *epochs* = 100
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers: $l = 0.0006$
- L2 regularization for the dense layer: $l = 0.01$

Results:

- train set accuracy: 0.805
- train loss: 1.07
- validation set accuracy: 0.81
- validation set loss: 1.08
- test set accuracy: **0.81**
- test set loss: 1.08




Figure 2: Model accuracy for LeNet-5 architecture - training & validation sets

Attempt 2

We used optuna to find hyperparameters for regularization, introduced decaying learning rate and increased number of epochs to 1000.

Hyperparameters & methods used:

- *optimizer*: Adam, with starting *learning rate* = 0.001, and decreasing by 50% each 150 epochs
- *epochs* = 1000
- *batch size* = 200
- *validation set size* = 20% of the training set
- *shuffle* the training data before each epoch
- L2 regularization for convolutional layers: $l \approx 0.00061$
- L2 regularization for the dense layer: $l \approx 0.0301$

Results:

- train set accuracy: 0.981
- train loss: 0.337
- validation set accuracy: 0.878
- validation set loss: 0.651
- test set accuracy: **0.892**
- test set loss: 0.629




Figure 3: Model accuracy for LeNet-5 with optimized hyperparameters

Conclusions:

The model after 20 epochs starts overfitting and validation set's accuracy improves very slowly.
(to-do)

5 Final model

As our final model we decided to have Lenet-5 because ...